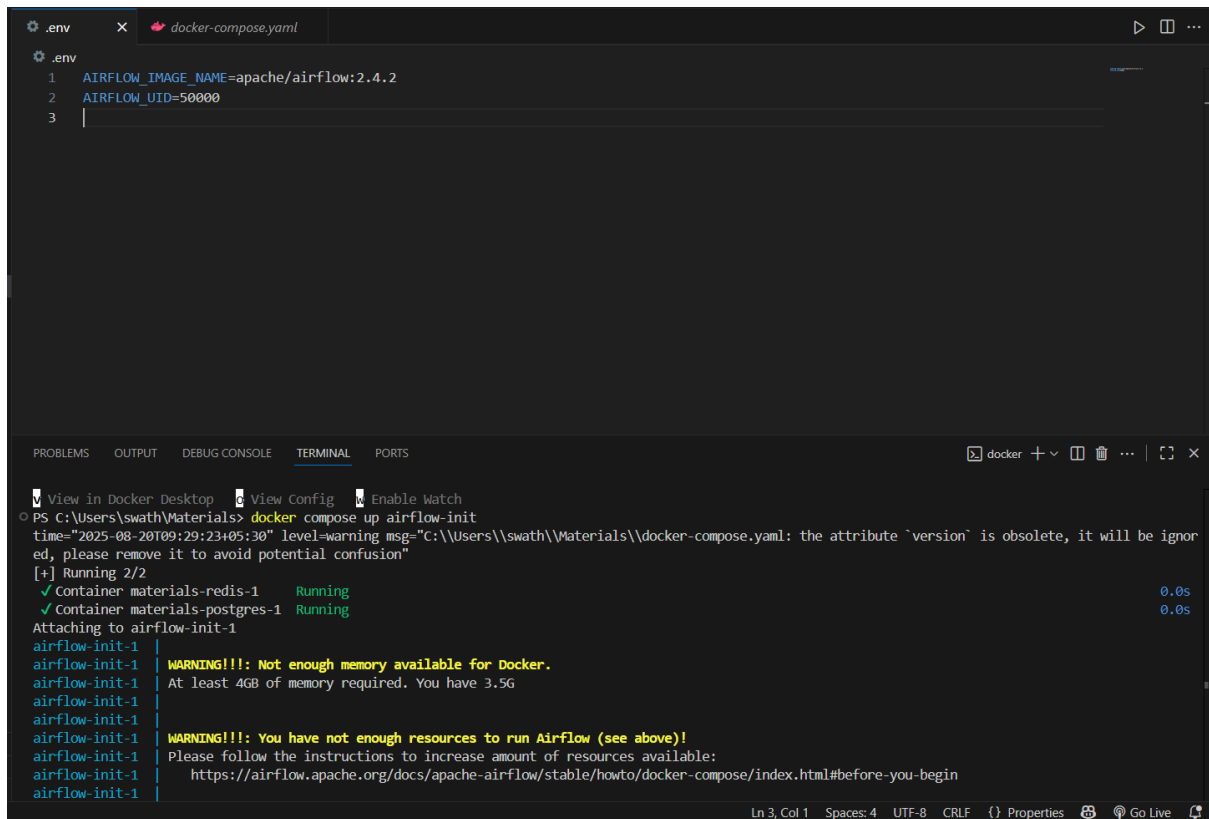


## ASSIGNMENT – 3 (Apache Airflow)

TRAINEE NAME: Swathi Baskaran

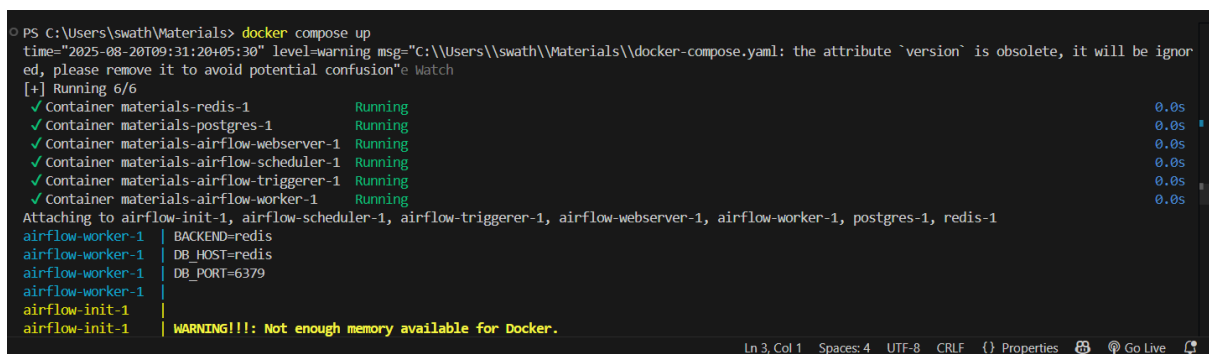
### Setting up airflow-init in VS Code



```
.env
1 AIRFLOW_IMAGE_NAME=apache/airflow:2.4.2
2 AIRFLOW_UID=50000
3 |

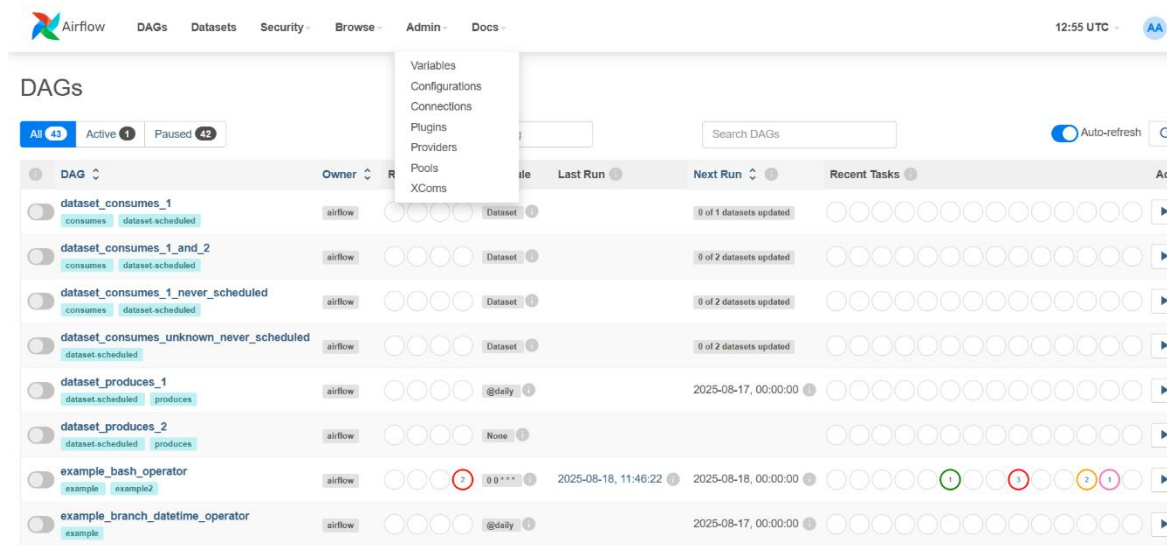
PS C:\Users\swath\Materials> docker compose up airflow-init
time="2025-08-20T09:29:23+05:30" level=warning msg="C:\\Users\\swath\\Materials\\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 2/2
 ✓ Container materials-redis-1      Running      0.0s
 ✓ Container materials-postgres-1   Running      0.0s
Attaching to airflow-init-1
airflow-init-1 |
airflow-init-1 | WARNING!!!: Not enough memory available for Docker.
airflow-init-1 | At least 4GB of memory required. You have 3.5G
airflow-init-1 |
airflow-init-1 |
airflow-init-1 | WARNING!!!: You have not enough resources to run Airflow (see above)!
airflow-init-1 | Please follow the instructions to increase amount of resources available:
airflow-init-1 | https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html#before-you-begin
airflow-init-1 |
```

### Starting up all services



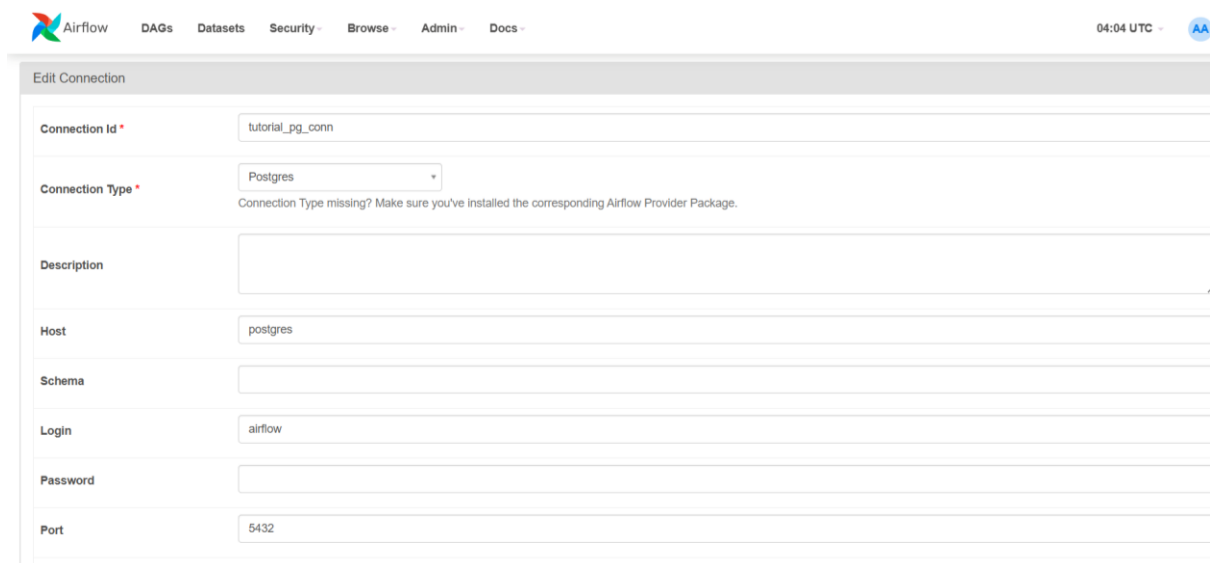
```
PS C:\Users\swath\Materials> docker compose up
time="2025-08-20T09:31:20+05:30" level=warning msg="C:\\Users\\swath\\Materials\\docker-compose.yml: the attribute `version` is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 6/6
 ✓ Container materials-redis-1      Running      0.0s
 ✓ Container materials-postgres-1   Running      0.0s
 ✓ Container materials-airflow-webserver-1 Running      0.0s
 ✓ Container materials-airflow-scheduler-1 Running      0.0s
 ✓ Container materials-airflow-triggerer-1 Running      0.0s
 ✓ Container materials-airflow-worker-1 Running      0.0s
Attaching to airflow-init-1, airflow-scheduler-1, airflow-triggerer-1, airflow-webserver-1, airflow-worker-1, postgres-1, redis-1
airflow-worker-1 | BACKEND=redis
airflow-worker-1 | DB_HOST=redis
airflow-worker-1 | DB_PORT=6379
airflow-worker-1 |
airflow-init-1 |
airflow-init-1 | WARNING!!!: Not enough memory available for Docker.
```

## Setting up a new connection



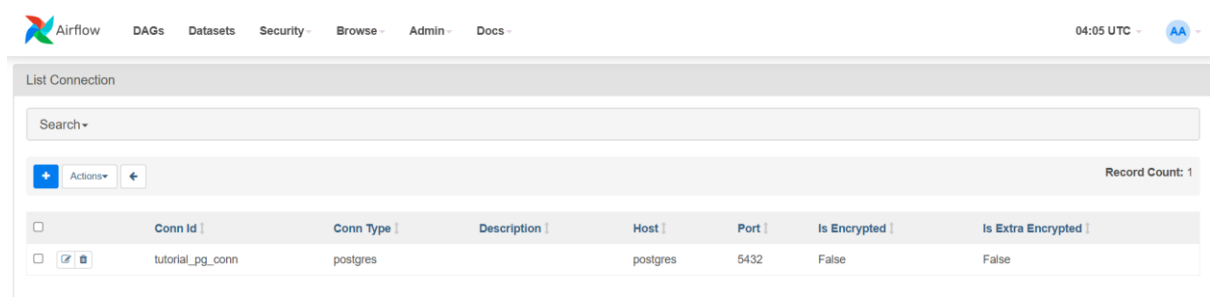
The screenshot shows the Airflow web interface. The top navigation bar includes links for DAGs, Datasets, Security, Browse, Admin, and Docs. The main header shows the time as 12:55 UTC. The 'DAGs' section is active, displaying a list of DAGs. A dropdown menu is open, showing options: Variables, Configurations, Connections, Plugins, Providers, Pools, and XComs. The DAG list includes entries like 'dataset\_consumes\_1', 'dataset\_consumes\_1\_and\_2', 'dataset\_consumes\_1\_never\_scheduled', 'dataset\_consumes\_unknown\_never\_scheduled', 'dataset\_produces\_1', 'dataset\_produces\_2', 'example\_bash\_operator', and 'example\_branch\_datetime\_operator'. Each entry shows its status (e.g., '0 of 1 datasets updated') and a 'Recent Tasks' section.

## Entering the connection details



The screenshot shows the 'Edit Connection' form in the Airflow web interface. The form is titled 'tutorial\_pg\_conn'. The 'Connection Type' is set to 'Postgres'. The 'Host' is 'postgres', 'Schema' is empty, 'Login' is 'airflow', 'Password' is empty, and 'Port' is '5432'. A message below the 'Connection Type' field states: 'Connection Type missing? Make sure you've installed the corresponding Airflow Provider Package.'

## Connection successfully created



The screenshot shows the 'List Connection' page in the Airflow web interface. The page displays a table with one connection listed. The table has columns for Conn Id, Conn Type, Description, Host, Port, Is Encrypted, and Is Extra Encrypted. The connection 'tutorial\_pg\_conn' is listed with a status of 'postgres' and '5432'.

Conn Id	Conn Type	Description	Host	Port	Is Encrypted	Is Extra Encrypted
tutorial_pg_conn	postgres		postgres	5432	False	False

## Create tables for staging and final data

```
.env employee_pipeline.py 1 X
dags > employee_pipeline.py > ...
1 from airflow.providers.postgres.operators.postgres import PostgresOperator
2
3 create_employees_table = PostgresOperator(
4     task_id="create_employees_table",
5     postgres_conn_id="tutorial_pg_conn",
6     sql="""
7         CREATE TABLE IF NOT EXISTS employees (
8             "Serial Number" NUMERIC PRIMARY KEY,
9             "Company Name" TEXT,
10            "Employee Markme" TEXT,
11            "Description" TEXT,
12            "Leave" INTEGER
13        );"""
14 )
15
16 create_employees_temp_table = PostgresOperator(
17     task_id="create_employees_temp_table",
18     postgres_conn_id="tutorial_pg_conn",
19     sql="""
20         DROP TABLE IF EXISTS employees_temp;
21         CREATE TABLE employees_temp (
22             "Serial Number" NUMERIC PRIMARY KEY,
23             "Company Name" TEXT,
24             "Employee Markme" TEXT,
25             "Description" TEXT,
26             "Leave" INTEGER
27         );"""
28 )
```

## Load data into the staging table

```
.env load_data_dag.py 2 X
dags > load_data_dag.py > get_data
1 import os
2 import requests
3 from airflow.decorators import task
4 from airflow.providers.postgres.hooks.postgres import PostgresHook
5
6
7 @task
8 def get_data():
9     # NOTE: configure this as appropriate for your airflow environment
10    data_path = "/opt/airflow/dags/files/employees.csv"
11    os.makedirs(os.path.dirname(data_path), exist_ok=True)
12
13    url = "https://raw.githubusercontent.com/apache/airflow/main/airflow-core/docs/tutorial/pipeline_example.csv"
14
15    response = requests.request("GET", url)
16
17    with open(data_path, "w") as file:
18        file.write(response.text)
19
20    postgres_hook = PostgresHook(postgres_conn_id="tutorial_pg_conn")
21    conn = postgres_hook.get_conn()
22    cur = conn.cursor()
23    with open(data_path, "r") as file:
24        cur.copy_expert(
25            "COPY employees_temp FROM STDIN WITH CSV HEADER DELIMITER AS ',' QUOTE '\"'",
26            file,
27        )
28    conn.commit()
```

## Merge and clean the data

```
.env merge_task.py x
dags > merge_task.py > merge_data
1 from airflow.decorators import task
2 from airflow.providers.postgres.hooks.postgres import PostgresHook
3
4
5 @task
6 def merge_data():
7     query = """
8         INSERT INTO employees
9         SELECT *
10        FROM (
11            SELECT DISTINCT *
12            FROM employees_temp
13        ) t
14        ON CONFLICT ("Serial Number") DO UPDATE
15        SET
16            "Employee Markme" = excluded."Employee Markme",
17            "Description" = excluded."Description",
18            "Leave" = excluded."Leave";
19    """
20    try:
21        postgres_hook = PostgresHook(postgres_conn_id="tutorial_pg_conn")
22        conn = postgres_hook.get_conn()
23        cur = conn.cursor()
24        cur.execute(query)
25        conn.commit()
26        return 0
27    except Exception as e:
28        return 1
```

## Defining the DAG

```
.env process_employees.py 4 x
dags > process_employees.py > ProcessEmployees
1 import datetime
2 import pendulum
3 import os
4 import requests
5
6 from airflow.decorators import dag, task
7 from airflow.providers.postgres.hooks.postgres import PostgresHook
8 from airflow.providers.postgres.operators.postgres import PostgresOperator
9
10
11
12 @dag(
13     dag_id="process_employees",
14     schedule="0 0 * * *", # daily at midnight
15     start_date=pendulum.datetime(2021, 1, 1, tz="UTC"),
16     catchup=False,
17     dagrun_timeout=datetime.timedelta(minutes=60),
18 )
19 def ProcessEmployees():
20     # Step 1: Create main employees table
21     create_employees_table = PostgresOperator(
22         task_id="create_employees_table",
23         postgres_conn_id="tutorial_pg_conn",
24         sql="""
25             CREATE TABLE IF NOT EXISTS employees (
26                 "Serial Number" NUMERIC PRIMARY KEY,
27                 "Company Name" TEXT,
28                 "Employee Markme" TEXT,
29                 "Description" TEXT,
30                 "Leave" INTEGER
31             );"""
32     )
33
34     # Step 2: Create staging table
35     create_employees_temp_table = PostgresOperator(
36         task_id="create_employees_temp_table",
37         postgres_conn_id="tutorial_pg_conn",
```

Ln 19, Col 24 Spaces: 4 UTF-8

```
1 import datetime
2 import pendulum
3 import os
4 import requests
5
6 from airflow.decorators import dag, task
7 from airflow.providers.postgres.hooks.postgres import PostgresHook
8 from airflow.providers.postgres.operators.postgres import PostgresOperator
9
10
11
12 @dag(
13     dag_id="process_employees",
14     schedule="0 0 * * *", # daily at midnight
15     start_date=pendulum.datetime(2021, 1, 1, tz="UTC"),
16     catchup=False,
17     dagrun_timeout=datetime.timedelta(minutes=60),
18 )
19 def process_employees():
20     get_data = task(...)
21     merge_data = task(...)
22     process_employees = task(...)
```

airflow-scheduler-1 | [2025-08-19 06:23:59,695] [scheduler\_job.py:537] INFO - Sending TaskInstanceKey(dag\_id='process\_employees', task\_id='merge\_data', run\_id='scheduled\_2025-08-18T00:00:00+00:00', try\_number=1, map\_index=1) to executor with priority 1 and queue default

airflow-scheduler-1 | [2025-08-19 06:23:59,695] [base\_executor.py:95] INFO - Adding to queue: ['airflow', 'tasks', 'run', 'process\_employees', 'merge\_data', 'scheduled\_2025-08-18T00:00:00+00:00', '--local', '--subdir', 'DAGS\_FOLDER/process\_employees.py']

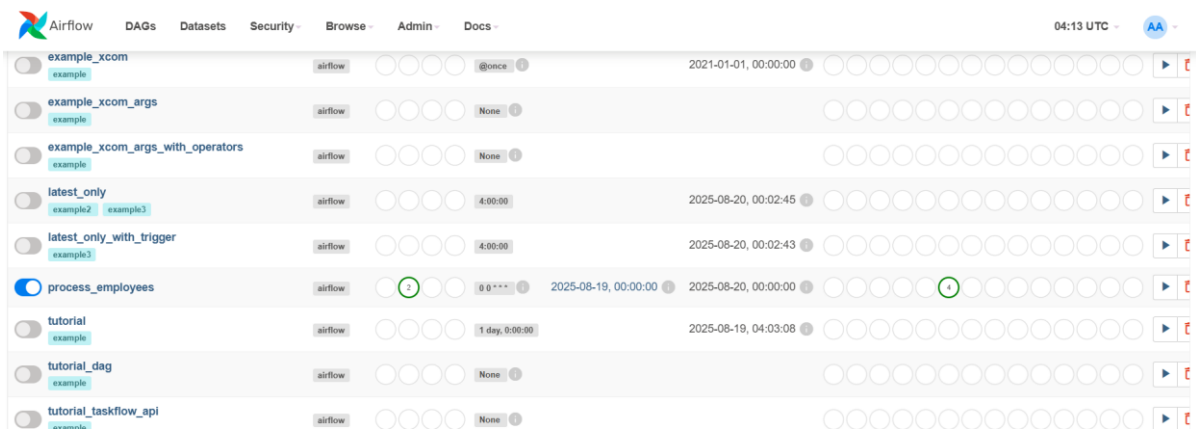
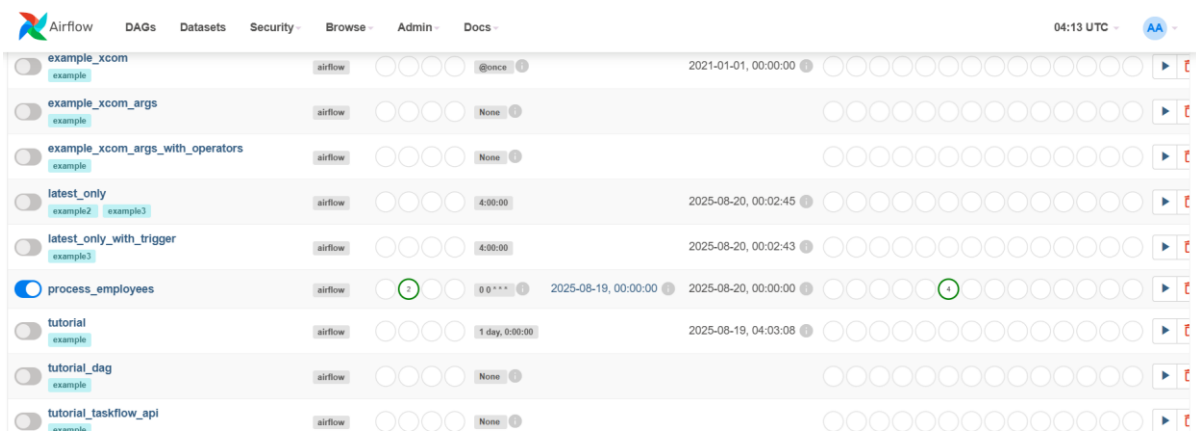
airflow-scheduler-1 | [2025-08-19 06:23:59,740] [scheduler\_job.py:595] INFO - Executor reports execution of process\_employees.merge\_data run\_id=scheduled\_2025-08-18T00:00:00+00:00 exited with status queued for try\_number 1

airflow-scheduler-1 | [2025-08-19 06:23:59,741] [scheduler\_job.py:595] INFO - Executor reports execution of process\_employees.get\_data run\_id=scheduled\_2025-08-18T00:00:00+00:00 exited with status success for try\_number 1

airflow-worker-1 | [2025-08-19 06:23:59,742] [INFO:forkpoolmarker-15] Filling up the DagBag from /opt/airflow/dags/process\_employees.py

airflow-scheduler-1 | [2025-08-19 06:23:59,748] [scheduler\_job.py:652] INFO - TaskInstance Finished: dag\_id=process\_employees, task\_id=get\_data, run\_id=scheduled\_2025-08-18T00:00:00+00:00, map\_index=1, run\_start\_date=2025-08-19 06:23:58.333827400+00:00, run\_end\_date=2025-08-19 06:23:59.550409400+00:00, run\_duration=1.216682000s

## Finding the DAG in the Apache UI



Airflow

DAGs

Datasets

Security

Browse

Admin

Docs

04:16 UTC

AA

DAG: process\_employees

Schedule: 0 0 \* \* \*

Next Run: 2025-08-20, 00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

<> Code

Audit Log

20-08-2025 04:15:51 AM

25

All Run Types

All Run States

Clear Filters

deferred failed queued removed restarting running scheduled shutdown skipped success up\_for\_reschedule up\_for\_retry upstream\_failed no\_status

Auto-refresh

Duration

00:03:19

00:01:39

00:00:00

create\_employees\_table

create\_employees\_temp\_table

get\_data

merge\_data

DAG Details

DAG Runs Summary

Total Runs Displayed	2
Total success	2
First Run Start	2025-08-19, 06:23:52 UTC
Last Run Start	2025-08-20, 03:59:29 UTC
Max Run Duration	00:03:19
Mean Run Duration	00:01:44

Airflow

DAGs

Datasets

Security

Browse

Admin

Docs

06:25 UTC

AA

DAG: process\_employees

success

Schedule: 0 0 \* \* \*

Next Run: 2025-08-19, 00:00

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

<> Code

Audit Log

2025-08-18T00:00:01Z

Runs

25

Run

scheduled\_\_2025-08-18T00:00:00+00:00

Layout

Left &gt; Right

Update

Find Task...

PostgresOperator PythonDecoratedOperator

deferred failed queued removed restarting running scheduled shutdown skipped success up\_for\_reschedule up\_for\_retry upstream\_failed no\_status

Auto-refresh

create\_employees\_table

create\_employees\_temp\_table

get\_data

merge\_data